

NexaMart Enterprise Data Warehouse — Milestone 2 Report

Group Arif · Arif (Lead) · Khallad (M2) · Sameen Fatima (M3) · Asad (M4) · Hamza (M5) · Azhar Sajjad (M6)

Team Contributions

Member	Deliverables
Arif (Lead)	Anomaly resolution (\$1), Gold rebuild (\$2), Validation suite (\$3), KPI views & reconciliation (\$4), Campaign conclusion (\$5), warehouse infrastructure & report assembly
Hamza, Asad, Khallad	Stakeholder dashboard — all 5 pages
Sameen Fatima, Azhar Sajjad	Executive presentation slides

Executive Summary (Owner: Arif, Lead)

Was the Back-to-School campaign (8–28 Aug 2024) actually successful, or did every team just count something different?

On a certainty-segregated, anomaly-corrected basis it was clearly positive. Campaign confirmed revenue ran about 50% above the pre-ramp run-rate (+₹1.27 Cr EC incremental). That is neither a failure nor the inflated +34% Sales celebrated, and the gap between those two readings is not a matter of opinion. It comes down to 24 resolved data anomalies. Sales had counted cancelled orders (A1, ₹4.08M) and seller-marked-sold listings (A4, ₹1.72 Cr raw value) as confirmed revenue, while a payment-after-cancel reversal liability (A2, flagged) and tax-basis mismatches (A3) distorted the rest. The brief's "+11%" reference figure is scenario framing (see S5); this warehouse does not reproduce it as a growth rate.

Getting to one number meant resolving every anomaly in corrected Silver with the audit trail kept and nothing deleted, rebuilding the affected Gold facts, validating with 10 checks, and exposing certainty-segregated KPI views in `NEXAMART_MARTS`. Confirmed revenue and Estimated Classified GMV (NexaLocal offline, ±35% band) are reported separately and never added together. That habit is where the legacy teams most often double-counted.

The headline numbers (the GSV → NCR waterfall is in S4, the certainty-labelled metrics in S5): GSV ₹141.5 Cr, less cancelled EC revenue (₹4.08M) and POS line discounts (₹2.75 Cr), gives NCR ₹138.35 Cr, the one trusted top-line. Return refunds (₹0.65 Cr) and Estimated Classified GMV (₹1.26 Cr point, ±35% band, ESTIMATED) are held separately and never added in. Over the campaign window, EC confirmed revenue is + ₹1.27 Cr incremental (B1 session-bridge attribution, INFERRED).

Section 1 — Anomaly Resolution Report (Owner: Arif, Lead)

For **each Category A** anomaly: detection method, root cause, the specific PySpark fix logic, the post-fix verification result (**affected row count before → after**), and the business impact in financial/operational terms. For **each Category B**: precise description of the ambiguity, the chosen interpretation, the PySpark implementation, the business justification for choosing it over alternatives, **and the quantified impact had the alternative been chosen** (brief §8.2).

Cross-references: detection = `sql/anomaly_discovery.sql` ; resolution code = `notebooks/05_anomaly_resolution.ipynb` ; per-anomaly docs = `sql/anomaly_resolution.sql`.

Detection verified (run 14 Jun 2026 as `NEXAMART_ENGINEER` against live `NEXAMART_SILVER`). 18 of 24 anomalies matched the expected fingerprints EXACTLY: A1=94, A2=27, A4=449, A5=5, A6=8, A7=10, A8=21, A9=1, A10=12, A11=178, A12=3, A13=3 rings, A14=18 strict, A15=25, A16=7, B1=126, B3=175, B7=25. The remaining six are, by design, **not single-count SQL detections**: A3 (schema-wide tax-basis comparison), B4 (catalogue-match threshold), B6 (Estimated-GMV signal model), and B5/B8 (probabilistic identity-merge / seller-trust scores) are computed in the `05` PySpark resolution; B2 (cross-period partial refund, 1 row) is finalised there too. Schema notes baked into the predicates: member-built Silver tables use base column names (`created_at` , `snapshot_date`), not the lead-built `_parsed` variants; `captured_ts` is a `TIMESTAMP` while `status_ts_parsed` is epoch-NUMBER; A9 is the cross-source `silver_ts_seller_listings` catalogue mismatch.

Category A

ID	Detection (live Silver predicate)	Root cause	Fix (method)	Before → After	Business impact
A1	<code>order_status='CANCELLED' AND subtotal_excl_tax>0</code> /94	cancelled orders left revenue on the line	<code>ZEROED_CANCELLED_REVENUE</code>	94 → 0	₹4.08M (₹0.41 Cr) over-count removed
A2	captured PG txn (<code>captured_ts</code> not null) after earliest <code>CANCELLED</code> status	capture fired post-cancel	<code>FLAGGED_REVERSAL_REQUIRED</code>	27 → 0	27 captured txns held as reversal liability (excl. NCR)
A3	POS incl-tax vs EC excl-tax basis differ	mixed tax bases summed naively	<code>NORMALISED_TAX_EXCLUSIVE</code>	schema-wide	₹16.25 Cr comparable basis

ID	Detection (live Silver predicate)	Root cause	Fix (method)	Before → After	Business impact
A4	<code>event_type_code='SELLER_SOLD' ✓449</code>	seller-sold counted as confirmed revenue	RELABELLED_ESTIMATED_GMV	449 → 0 confirmed	₹1.72 Cr moved to ESTIMATED
A5	<code>atp_qty>0 AND physical_qty=0 ✓5</code>	ATP not decremented to physical	CORRECTED_ATP_TO_ZERO	5 → 0	oversell risk removed
A6	<code>sellable_qty<0 OR physical_qty<0 ✓8</code>	negative inventory written	CORRECTED_ATP_TO_ZERO	8 → 0	stockout analysis fixed
A7	<code>inspection_status='PENDING' AND restocked_qty>0 ✓10</code>	restock before inspection	ZEROED_RESTOCK_PRE_INSPECTION	10 → 0	ATP inflation removed
A8	stores 3/7/12 × (1–7 Aug) absent from snapshots	snapshots not captured in ramp-up	RECONSTRUCTED_SNAPSHOT	~21 reconstructed	allocation accuracy (1–7 Aug)
A9	<code>silver_ts_seller_listings</code> : listing 42 / NX-TECH-0001 catalogue mismatch	seller feed contradicts catalogue	APPLIED_CANONICAL_PRODUCT	1 → 0	category revenue corrected
A10	<code>condition_on_receipt='OPENED' AND restocked_as_condition='NEW'</code>	open-box restocked as NEW	CORRECTED_CONDITION_OPEN_BOX	12 → 0	price premium quantified
A11	<code>customer_id='9999' ✓178</code>	guest placeholder collides w/ real loyalty	REKEYED_GUEST_BUCKET	178 → 0	customer profile de-collided
A12	self-join <code>image_hash</code> +seller, prior SOLD/EXPIRED, ACTIVE relist	listing relisted after sold	LINKED_RELISTING_EXCLUDED	3 pairs	GMV double-count removed
A13	<code>image_hash</code> shared by ≥2 seller accounts ✓3 rings	coordinated listing ring	FLAGGED_FRAUD_RING	rings flagged	trust & safety
A14	<code>delivered_ts<pickup_ts</code> (strict) ✓18	courier clock drift	CORRECTED_DELIVERY_TS / ESCALATED_MANUAL_REVIEW	18 (>72h escalated)	on-time rate corrected
A15	<code>days_post_delivery<0</code>	review before delivery	SET_VERIFIED_PURCHASE_FALSE	25 → 0 verified	review-score KPI corrected
A16	<code>is_duplicate_flag=TRUE ✓7</code>	same complaint via many channels	DEDUPED_CANONICAL_CASE	7 → deduped	complaint volume corrected

Category B

ID	Ambiguity	Chosen interpretation	Quantified alternative
B1	promo-less in-window + prior UTM	attribute 102 (conf 0.85)	not attributing understates ~₹0.94 Cr confirmed
B2	Aug sale / Sep refund	recognise in return period	original-period reduces Aug NCR ~₹2.16 L
B3	175 PCK null-ref movements	probable missing-ref (INFERRED)	shrinkage reading overstates loss
B4	NL title → catalogue	≥0.75 match / 0.65–0.75 review	lower threshold injects false matches
B5	cross-channel identity	merge ≥0.90	0.70 over-merges distinct customers
B6	Estimated GMV formula	weighted signals + ±35% band	seller-sold-only mis-states
B7	BOPIS no pickup event	treat as fulfilled (scan miss)	not-collected understates fulfilment
B8	seller trust score	weighted 8-signal → 5 tiers	equal-weight unacceptable

Per-anomaly detail (detection → root cause → fix → verification → impact)

Resolution code: `notebooks/05_anomaly_resolution.ipynb` (one cell per anomaly); post-resolution verify: `sql/anomaly_resolution.sql`. Every corrected row keeps `anomaly_flag=TRUE` + its original `anomaly_reason_code` and gains `resolution_applied=TRUE` + `resolution_method` (+ `b_classification` for Category B), and nothing is deleted (brief §8.1).

- A1 — Cancelled orders carrying revenue.** 94 EC orders had `order_status='CANCELLED'` but a positive `subtotal_excl_tax`; the cancellation workflow set the status without reversing the revenue line. Fix: flag `CANCELLED_WITH_REVENUE`, zero the revenue (`ZEROED_CANCELLED_REVENUE`). In Gold the keystone fact keeps the gross line value (for GSV) but sets `confirmed_revenue_excl_tax=0` for cancelled lines, so NCR excludes them. Verify → 0 unresolved. **Removes the ₹4.08M (₹0.41 Cr) cancelled-revenue over-count that inflated Sales' +34% claim** (`vw_revenue_leakage` CANCELLATION = ₹4,078,605).
- A2 — Payment captured after cancellation.** 27 captured PG transactions (26 strictly after the cancellation timestamp) sit against cancelled orders. Fix: `FLAGGED_REVERSAL_REQUIRED` (kept, excluded from NCR), held as a reversal liability. The captured-payment exposure sits in the PG source, outside the Gold facts, and is never counted as revenue.
- A3 — Mixed tax bases.** POS stores tax-inclusive amounts, EC tax-exclusive; naively summing them is meaningless. Fix: derive a tax-exclusive POS measure (`NORMALISED_TAX_EXCLUSIVE`) so every channel reconciles on one basis. Reconciles the ~₹16.25 Cr naive cross-channel figure.

- **A4 — Seller-marked-sold counted as revenue.** 449 `SELLER_SOLD` NexaLocal events were summed as confirmed GMV. Fix: relabel certainty `ESTIMATED` (`RELABELLED_ESTIMATED_GMV`); they feed the B6 model at weight 0.60, never Confirmed GMV. Moves ₹1.72 Cr out of confirmed totals.
- **A5 / A6 — Impossible inventory.** 5 warehouse SKU-dates showed `ATP>0` with `physical=0`; 8 store snapshots carried negative quantities. Both corrected to 0 (`CORRECTED_ATP_TO_ZERO` , A6 also sets an oversell flag), removing oversell risk and fixing stockout analysis.
- **A7 / A10 — Return-receipt handling.** 10 receipts restocked while `inspection_status='PENDING'` (`ZEROED_RESTOCK_PRE_INSPECTION` , ATP inflation removed); 12 open-box returns mislabelled `NEW` (`CORRECTED_CONDITION_OPEN_BOX` , price-premium quantified). Distinct predicates, overlapping sets.
- **A8 — Missing snapshot days.** Stores 3/7/12 each missing 7 snapshot days (1–7 Aug, pre-ramp). The reconstructed sibling `silver_store_inventory_snapshots_reconstructed` (last-known + interpolation) is resolution-stamped `RECONSTRUCTED_SNAPSHOT` at certainty `INFERRED`, never `CONFIRMED` and never inside the campaign window.
- **A9 — SKU/product mismatch.** Listing 42 cites SKU `NX-TECH-0001` (a laptop) but describes a phone case. Fix: catalogue wins (`APPLIED_CANONICAL_PRODUCT`), 1 → 0.
- **A11 — Placeholder-ID collision.** 178 guest EC orders share `customer_id='9999'` with a real dormant loyalty account (Sarah Chen). Fix: rekey guests to `GUEST-{session_id}` (`REKEYED_GUEST_BUCKET`), de-colliding the real customer's profile.
- **A12 / A13 — Classified integrity.** 3 relisted-after-sold pairs (same seller + image hash) linked and excluded from GMV (`LINKED_RELISTING_EXCLUDED` , reliability `LOW`); 3 image-hash rings spanning ≥ 2 seller accounts (18 listings) flagged `FLAGGED_FRAUD_RING` .
- **A14 — Courier clock drift.** 18 shipments show a `DELIVERED` event before `PICKED_UP`. For $|\text{delta}| \leq 72\text{h}$ the delivered timestamp is corrected to pickup + 36h (`CORRECTED_DELIVERY_TS`); beyond 72h it is escalated rather than auto-corrected (`ESCALATED_MANUAL_REVIEW`), which is the expected residual after resolution.
- **A15 / A16 — CX signals.** 25 reviews posted before delivery set `is_verified_purchase=FALSE` (`SET_VERIFIED_PURCHASE_FALSE`); 7 duplicate support cases resolved to their canonical key (`DEDUPED_CANONICAL_CASE`) so complaint volume counts distinct incidents.
- **B1 — Attribution bridge.** Of 126 promo-less in-window orders, 102 had a prior `BTS2024` UTM session within the 2h bridge → `ATTRIBUTED` at confidence 0.85. **Alternative:** not attributing understates campaign revenue by ~₹0.94 Cr confirmed (₹9,432,000 over the 102 attributed orders; ₹1.02 Cr gross).
- **B2 — Partial-refund period.** 1 cross-month partial refund recognised in the `return` period (`RECOGNISE_IN_RETURN_PERIOD`); both period impacts are exposed. Original-period recognition would cut Aug NCR by ~₹2.16 L.
- **B3 — Null-reference movements.** 175 `NULL- reference_number` picks (all PCK) classified `PROBABLE_MISSING_REF` (`INFERRED`, back-fill lag) rather than treated as shrinkage, which would overstate loss.
- **B4 — Catalogue matching.** NexaLocal free-text titles matched to the catalogue with `rapidfuzz`: ≥ 0.75 `MATCHED`, 0.65–0.75 `MANUAL_REVIEW`, < 0.65 `UNMATCHED`. A lower threshold injects mis-linked demand signals.
- **B5 — Identity merge.** A single 0.92-confidence cross-channel pair merged (`MERGED` , threshold ≥ 0.90); a 0.70 threshold would over-merge distinct customers.
- **B6 — Estimated Classified GMV.** Locked weights (`SELLER_SOLD` 0.60 / `PHN_REVEAL` 0.15 / `CHAT` 0.08 / `OFFER_ACC` 0.30) $\times$ listing confidence, $\pm 35\%$ band, labelled `ESTIMATED`. It is surfaced only via `vw_estimated_classified_gmv` and never summed into confirmed revenue.
- **B7 — BOPIS scan miss.** 25 BOPIS orders marked `DELIVERED` with no `BOPIS_COLLECTED` scan treated as fulfilled (`TREAT_AS_FULFILLED` , ~13% scan-miss baseline) + `collection_unconfirmed=TRUE` , excluded from the BOPIS SLA.
- **B8 — Seller trust.** Each seller carries a composite trust score built from eight behavioural signals. It is deliberately not an equal-weight average (the brief rules that out); the weights reflect how much each signal damages buyer trust. Order-cancellation rate and late-fulfilment rate carry the most (0.20 each), since a broken delivery promise is the most direct harm. Complaints per 100 orders (0.15) and the integrity signals come next: NexaLocal duplicate-listing rate (0.12), report rate (0.10) and moderation actions (0.08). Return rate (0.10) and buyer-contact response rate (inverted, 0.05) are lagging or service indicators. The weights sum to 1.00 and combine as $\text{trust_score} = 1 - \sum(w_i \times \text{normalised_risk}_i)$, so a clean seller scores near 1 and a problem seller scores low. The score maps to the five tiers in `dim_seller_risk_tier` : `VERIFIED_TRUSTED` (≥ 0.85), `STANDARD` (≥ 0.65), `UNDER_REVIEW` (≥ 0.40), `HIGH_RISK` (≥ 0.20) and `SUSPENDED` (below 0.20); a flagged seller escalates to `UNDER_REVIEW`, which is reversible 30-day monitoring rather than suspension. The deployed `vw_seller_risk_score_distribution` rolls these five tiers into a simpler `LOW` / `MEDIUM` / `HIGH` risk band for the dashboard, and on the current data all 348 active sellers sit in the `LOW`-risk (trusted) band, so `UNDER_REVIEW` and below are defined by the model but unpopulated for this period.

Section 2 — Gold Rebuild Summary (Owner: Arif, Lead)

Corrected Silver flows to Gold through partial rebuilds (`notebooks/06_gold_rebuild.ipynb`): only the facts an affected Silver entity feeds are re-pointed at corrected Silver and overwritten. The 8 static lookup dims (`dim_date` , `dim_channel` , `dim_step` , `dim_payment_method` , `dim_delivery_method` , `dim_listing_condition` , `dim_return_reason` , `dim_promotion`) have no Silver feed, so they are rebuild-exempt.

Facts rebuilt (9): `fact_ecommerce_order_line` (A1/A2/A3/A11/B1 — keystone, with the GSV/NCR split), `fact_order_fulfilment` (A14/B7), `fact_store_inventory_snapshot` (A6/A8), `fact_warehouse_inventory_snapshot` (A5), `fact_inventory_transaction` (B3), `fact_classified_listing_snapshot / _event` (A4/A12/A13/B6), `fact_customer_review` (A15), `fact_customer_complaint` (A16). Each rebuild carries an inline `assert_grain` (`COUNT(*) = COUNT(DISTINCT grain)`).

Fact grain declarations. Every fact is built and validated at one declared grain (validation Check 3 enforces `COUNT(*) = COUNT(DISTINCT grain)`). Keys shown in code font are the exact declared keys; the rest are stated by grain.

Fact	Grain — one row per	Declared key
fact_ecommerce_order_line	EC order line	order_id, line_id
fact_store_sale_line	store POS sale line	transaction × line
fact_return_line	return line	return × line
fact_order_fulfilment	EC order (fulfilment lifecycle)	order_id
fact_store_inventory_snapshot	store × product × day (semi-additive)	store_id, product_code, snapshot_date
fact_warehouse_inventory_snapshot	SKU × location × day (semi-additive)	sku × location × date
fact_inventory_transaction	inventory movement	movement_id, node_type
fact_web_session	web session (aggregate clickstream grain)	session
fact_web_page_event	page event (atomic clickstream grain, Kimball Ch. 15)	session_id, event_id
fact_classified_listing_snapshot	listing snapshot (semi-additive)	listing_id
fact_classified_listing_event	NexaLocal listing event	listing event
fact_seller_performance_snapshot	seller × period (semi-additive)	seller × period
fact_customer_review	product/order review	review
fact_customer_complaint	support case	case_id

The two clickstream grains (session aggregate plus atomic page event) follow Kimball Chapter 15; the snapshot facts are semi-additive and are never summed across dates in the KPI views (validation Check 5).

Row counts. Corrections change measures and flags, not grain, so before/after row counts are identical per table. The notebook's final cell prints the before/after delta as proof, and a non-zero delta is treated as a defect. The descriptive dims (`dim_customer / dim_product / dim_seller`) are left intact: facts carry their dim FK as `surrogate_key(natural_code)`, so every fact→dim join still resolves, and the SCD2 dim refresh that carries the A9/B5/B8 descriptive corrections is a member follow-up. `fact_store_sale_line`, `fact_return_line` and `fact_seller_performance_snapshot` have no in-repo M1 build logic, so they are noted here rather than rebuilt blind.

Before → after row counts (Phase-B run, 15 Jun 2026; every delta 0, confirming grain-preserving corrections):

Rebuilt fact	Before	After	Δ
fact_ecommerce_order_line	1,840	1,840	0
fact_order_fulfilment	963	963	0
fact_store_inventory_snapshot	217,800	217,800	0
fact_warehouse_inventory_snapshot	38,610	38,610	0
fact_inventory_transaction	468,455	468,455	0
fact_classified_listing_snapshot	1,253	1,253	0
fact_classified_listing_event	38,706	38,706	0
fact_customer_review	377	377	0
fact_customer_complaint	129	129	0

All nine rebuilt facts pass their inline `assert_grain ( COUNT(*) = COUNT(DISTINCT grain) )`, and a second full Run-all reproduces these counts identically (idempotent). The zero deltas are the correctness proof: A1, A2, A4, A14, B1, B3 and the rest re-value measures and set flags on existing rows, so no row is added or deleted.

Section 3 — Validation Outcomes (Owner: Arif, Lead)

`sql/validation_suite.sql` runs 10 checks against `NEXAMART_GOLD` after every rebuild; each returns the offending rows (target 0, except Check 7 which expects ≥1 row per fact in the campaign window).

#	Check	What it proves
1	Completeness	all 27 Gold tables non-empty (UNION-ALL row counts)
2	Referential integrity	no orphan <code>date_key</code> (computed identically everywhere) across the facts
3	Grain	<code>COUNT(*) = COUNT(DISTINCT grain)</code> per declared grain
4	Additive sanity	<code>net = gross - discount</code> on <code>fact_store_sale_line</code> (no return column at this grain; returns live in <code>fact_return_line</code>)
5	Semi-additive guard	no MARTS view <code>SUM</code> s ATP across dates (<code>vw_atp_sku_loc_date</code> exposes the grain, no <code>SUM</code>)

#	Check	What it proves
6	Certainty completeness	no NULL <code>metric_certainty_level</code> in any of the 14 facts
7	Campaign coverage	≥ 1 row per campaign-accepting fact in 8–28 Aug (via <code>dim_date</code> join)
8	Inventory reconciliation	snapshot $\Delta(\text{physical_qty}) = \text{net signed movement, per store}\times\text{product}$
9	Certainty segregation	no Finance view row is ESTIMATED while <code>is_confirmed_transaction</code> <> FALSE ; the estimated view is 100% ESTIMATED
10	Temporal consistency	no auto-correctable (≤72h) delivered-before-picked remains post-A14 (>72h rows flagged UNRELIABLE)

Iteration log (from the Phase-B runs, 15 Jun 2026; honest failure reporting is graded above an implausible first-run pass):

Iteration	Checks failed	Root cause	Fix applied	Re-run result
1	Check 4 (additive sanity); Check 8 (inventory reconciliation)	C4 referenced <code>return_amount</code> , which does not exist on the member-built <code>fact_store_sale_line</code> (real measures: <code>gross_amount</code> , <code>discount_amount</code> , <code>net_amount</code> , <code>tax_amount</code> , <code>cogs_amount</code>) → invalid-identifier error. C8 <code>moves</code> CTE summed all-time store movements against a campaign-window snapshot delta — incomparable by construction (1,095 offending pairs).	C4 → assert the identity that actually holds at this grain, <code>net_amount = gross_amount - discount_amount</code> . C8 → window <code>moves</code> to 8–28 Aug to match the snapshot bounds (<code>quantity_delta</code> is already signed: PICK/DMG –, RCVD/RET +).	C4 = 0 broken of 24,507 (PASS). C8 = 863 of 1,100 store×product pairs still residual (237 reconcile exactly).
2 (final)	none blocking	C8 residual root-caused: store inventory snapshots are independent periodic physical counts , while <code>fact_inventory_transaction</code> records only PICK/RCVD/RET/DMG — store POS sell-through is not in the movement ledger , so the two facts cannot reconcile to unit tolerance (avg residual 7.4 units; 515 pairs within ±2). This is a data-sourcing observation, not an anomaly-resolution defect; tolerance 0 is left deliberately strict so the check keeps surfacing the gap.	Documented (no code change).	9 of 10 checks PASS ; Check 8 = documented reconciliation residual.

Final state per check (live NEXAMART_GOLD / NEXAMART_MARTS, run 15 Jun 2026): C1 27/27 tables non-empty (min row_count 2); C2 0 orphan `date_key`; C3 0 duplicate-grain groups; C4 0 broken of 24,507; C5 0 ATP-summing views; C6 0 NULL `metric_certainty_level`; C7 every campaign-accepting fact ≥ 1 in window (min 227); C8 documented residual (above); C9 0 certainty-segregation breaches; C10 0 auto-correctable delivered-before-picked remaining.

Section 4 — KPI Reconciliation Report (Owner: Arif, Lead)

The seven teams disagreed because each summed a different measure on a different basis. The one trusted number, Net Confirmed Revenue (NCR), is reached from Gross Sale Value (GSV) through a named, quantified waterfall (`NEXAMART_MARTS.vw_gsv`, `vw_ncr`, `vw_revenue_leakage`).

GSV → NCR waterfall (real figures, deployed views, run 15 Jun 2026)

The deployed `vw_gsv` / `vw_ncr` quantify two GSV→NCR deductions: A1 cancelled EC revenue, and POS line-level discounts (store NCR = `net_amount = gross_amount - discount_amount`). A2 (payment-after-cancel) carries no extra ₹, because those payments sit on orders already zeroed by A1 and are held as a reversal *liability* rather than revenue. A3 normalises POS to a tax-exclusive basis, which is a basis correction rather than a ₹ line, and B2's single partial refund (₹2.16 L) is recognised in the return period. All figures below span the four in-scope phases (1 Jan – 14 Sep 2024).

Gross Sale Value (GSV, tax-exclusive, all confirmed channels)	₹141.51 Cr	(₹1,415,087,895)
- Cancelled EC revenue	← A1 ZEROED_CANCELLED_REVENUE	₹0.41 Cr (₹4,078,605; 94 orders / 157 lines)
- POS line-level discounts	← gross→net at the store grain	₹2.75 Cr (₹27,502,800)
= Net Confirmed Revenue (NCR)	₹138.35 Cr	(₹1,383,506,490) ← the trusted top-line

Two further measures are surfaced separately and never folded into NCR. Platform return refunds (`vw_revenue_leakage`, `RETURN_REFUND`) come to ₹0.65 Cr (₹6,450,710, from `fact_return_line.refund_amount`), and Estimated Classified GMV (NexaLocal offline, `vw_estimated_classified_gmv`, ESTIMATED, ±35% band) is ₹1.26 Cr at the point estimate (₹12,603,544; band ₹0.82 Cr – ₹1.70 Cr). That estimate is B6-weighted (SELLER_SOLD 0.60 / PHN_REVEAL 0.15 / CHAT 0.08 / OFFER_ACC 0.30 × listing `asking_price`); the raw seller-marked-sold value A4 pulled out of confirmed GMV is ₹1.72 Cr (`SUM(asking_price)` over the 449 SELLER_SOLD listings). Confirmed GMV (`vw_confirmed_gmv`, equal to NCR by construction) and Estimated Classified GMV are reported separately and never added, which was the single most common legacy double-count.

Data note: `cogs_amount` is unpopulated in the member-built sale facts, so `vw_gross_margin_by_channel` / `vw_net_margin_after_fulfilment` currently return margin = revenue (numerator = denominator); margin % is therefore not yet meaningful and is flagged for the SCD/COGS member follow-up rather than reported as a KPI.

Why each team's number differed — claimed vs warehouse-reproducible

Each team's headline came from summing a different measure on a different basis. The brief's claimed growth rates (Sales +34%, Finance +11%, ...) are **not reproducible as growth rates** from this warehouse; what the deployed views *do* reproduce is the ₹ figure each team mis-stated and the corrected one beside it. Every figure below traces to a live NEXAMART_MARTS view (run 15 Jun 2026) and carries its certainty.

Team	Claimed (brief)	Warehouse-reproducible figure (certainty)	Anomaly(ies)
Sales	+34% revenue	confirmed revenue overstated by ₹2.13 Cr ; -₹0.41 Cr cancelled EC orders removed (vw_revenue_leakage , CONFIRMED) + ₹1.72 Cr seller-marked-sold reclassified out of confirmed to ESTIMATED	A1, A4/B6
Finance	+11% (reference)	NCR ₹138.35 Cr — the trusted basis every other number reconciles to (CONFIRMED)	—
Marketplace	inflated GMV	Estimated Classified GMV ₹1.26 Cr B6-weighted vs ₹1.72 Cr raw (-₹0.46 Cr over-weight; ESTIMATED) + duplicate-listing inflation 1.015× (19 of 1,253 listings; CONFIRMED)	B6, A12
Inventory	"only 3 stockouts"	5 SKU-location-days at ATP = 0 and 5 residual oversell-risk snapshots, all inside the campaign window ; inventory accuracy 99.47% (216,645/217,800; CONFIRMED)	A5, A6
Ecommerce	abandonment +22%	cart abandonment 41.5% (684/1,647), payment-failure 7.0% (67/963), on-time delivery only 54.2% (340/627) — abandonment tracks supply/delivery failure, not demand (CONFIRMED, whole-period)	A5/A6, A14
Store Ops	pickup handling ×3	896 BOPIS orders made visible (avg readiness 16.9 h); BORIS = 0 in window (CONFIRMED)	B7, BORIS split
Support	340 complaints	A16 cross-channel case de-duplication applied in S1; complaint volume is not surfaced as a MARTS KPI (no support view)	A16
Cross-cutting	—	tax-basis (A3) normalises POS to tax-exclusive; refund-period (B2) moves the ₹2.16 L partial refund to the return period — both make cross-channel sums comparable	A3, B2

Grain note: the funnel/conversion and seller figures above are **whole-period** (those views are not phase-split); the phase-split numbers (GSV / NCR / Confirmed GMV by Baseline | Campaign | Post) are in the waterfall above and in S5.

On the campaign window itself (8–28 Aug, 21 days), vw_campaign_incremental_revenue puts EC confirmed campaign revenue at **₹1.79 Cr** against a length-normalised baseline of **₹0.52 Cr → +₹1.27 Cr incremental (CONFIRMED)**; B1's session-bridge attribution (102 of the 126 promo-less in-window orders, confidence 0.85) is the INFERRED mechanism behind that lift. The data-backed campaign verdict is in S5.

Section 5 — Campaign Performance Conclusion (Owner: Arif, Lead)

On a certainty-segregated, anomaly-corrected basis the Back-to-School campaign was clearly positive. Campaign-window confirmed revenue ran at ₹11.69M/day against the ₹7.78M/day pre-ramp run-rate, a +50% lift, and it sat well above the longer baseline too. In absolute terms that is +₹1.27 Cr of EC incremental confirmed revenue (vw_campaign_incremental_revenue : campaign ₹1.79 Cr vs a length-normalised baseline ₹0.52 Cr). The legacy +34% Sales headline was inflated by counting two things that are not confirmed revenue: cancelled EC orders (A1, ₹4.08M, now zeroed) and seller-marked-sold classified listings (A4, ₹1.72 Cr raw value, now relabelled ESTIMATED and reported only in the ±35% band, with a B6 point estimate of ₹1.26 Cr). Stripping out that inflation is not a judgement call; it follows directly from the 24 resolved anomalies. The brief's specific "+11% vs +34%" figures are scenario framing and do not reproduce as growth rates from this warehouse, so the defensible verdict is the clean positive uplift above, with the inflation kept separate rather than collapsed into one contested percentage. Two operational caveats remain: on-time delivery at 54% (vw_on_time_delivery_rate , 340/627) and the inventory snapshot-versus-ledger reconciliation residual (S3). Both are conditions the campaign exposed rather than caused.

Each conclusion metric is reported with its certainty level and never conflated; ESTIMATED values are never summed into CONFIRMED totals:

Conclusion metric	Source view	Certainty	Reading (live, 15 Jun 2026)
Net Confirmed Revenue	vw_ncr , vw_campaign_incremental_revenue	CONFIRMED	₹138.35 Cr total; campaign ₹24.55 Cr ; + ₹1.27 Cr EC incremental
Confirmed GMV	vw_confirmed_gmv	CONFIRMED	₹138.35 Cr (= NCR by construction); campaign-phase ₹24.55 Cr
Estimated Classified GMV (±35%)	vw_estimated_classified_gmv	ESTIMATED	₹1.26 Cr point (band ₹0.82–1.70 Cr); reported separately, never added to NCR
Net margin after fulfilment	vw_net_margin_after_fulfilment	CONFIRMED	deferred — cogs_amount unpopulated in source facts; margin % not yet meaningful (member COGS follow-up)
Inventory health (stockout / accuracy)	vw_stockout_rate , vw_inventory_accuracy_rate	CONFIRMED	accuracy 216,645/217,800 = 99.5% ; stockouts ≤ 1/195 SKU-day on only a few campaign days (A5/A6 corrected)
Fulfilment (on-time, BOPIS)	vw_on_time_delivery_rate , vw_bopis_pickup_readiness_time	CONFIRMED	on-time 340/627 = 54% ; BOPIS readiness 16.9h avg over 896 orders (A14 corrected; >72h escalated)
Customer satisfaction	vw_validated_report_rate , review/complaint views	CONFIRMED / INFERRED	validated-report 18/19 ; A15 (25 reviews) set unverified, A16 (7 cases) deduped
Classified engagement	vw_listing_contact_rate , vw_active_listing_count	CONFIRMED	538 active listings, 564 contacts; relisting 3/1,253; A12/A13 fraud rings excluded

Certainty discipline holds end-to-end: validation Check 9 confirms 0 ESTIMATED rows leak into any confirmed Finance view and `vw_estimated_classified_gmv` is 100% ESTIMATED; Check 5 confirms no view sums ATP across dates. The ₹1.26 Cr estimated classified GMV is never added to the ₹138.35 Cr NCR; the two sit in different certainty bands by design.